

# Data Augmentations for Arithmetic Length Generalization in Transformers

Lynnix Zou, Muhammad H. Ashiq, Grigoris Chrysos  
University of Wisconsin-Madison



## Introduction

Transformers achieve impressive performance on many tasks, yet they often fail to generalize to longer input lengths even for basic arithmetic such as multi-digit addition. Existing methods [1, 2] usually modify the architecture, for example by introducing new position markers to positional embeddings.

We ask whether such machinery is actually required once a standard Transformer already fits the training distribution. We study whether a plain decoder-only Transformer can achieve strong length generalization on arithmetic purely through a data format.

## Methods

A central difficulty in arithmetic length generalization is identifying which digits should be operated on when sequence length grows. Transformers tend to retrieve digits by absolute index, so a local update rule does not automatically extend to unseen positions. Prior work [3] shows that the model learns a per-digit rule of the form:

$$y_i = (a_i + b_i + c_{i-1}) \bmod 10, \quad c_i = \left\lfloor \frac{a_i + b_i + c_{i-1}}{10} \right\rfloor$$

Our idea is to remove this bottleneck by forcing corresponding digits in the operands and the sum to share the same relative positions inside a longer padded layout. The Transformer can then reuse the same per-digit computation across all positions that become visible through padding.

## Results

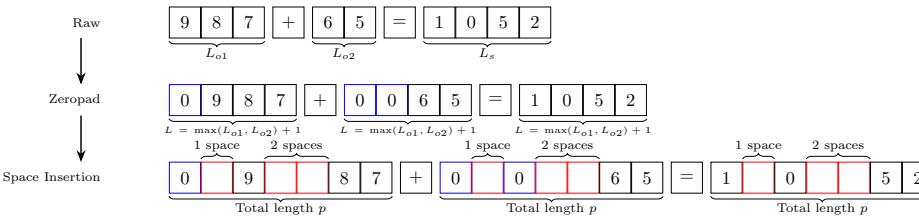


Figure 1: Illustration of the Aligned Blankspace Augmentation (ABA) method on addition

**ABA** is achieved by:

1. **Zero-Padding:** First, each numerical in the equations zero-padded to length  $L$ . This length  $L$  is determined to be sufficiently long to the expected maximum length of the sum.
2. **Synchronized Space Insertion:** An identical pattern of blank spaces is applied to expand the numbers to a target length,  $p$ . We first sample  $L$  unique positions from the  $p$  available slots uniformly at random without replacement, then insert The digits sequentially, with any unfilled positions becoming blank spaces.

## Results on addition

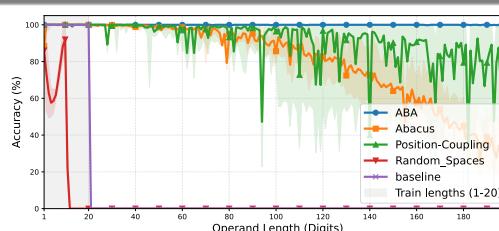


Figure 2: Methods for Length Generalization on Addition.

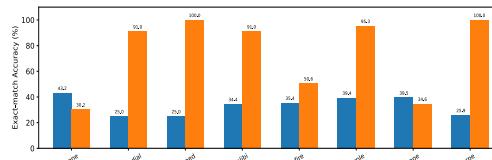


Figure 3: Different positional embeddings using ABA.

## Results

Task	baseline	ABA	Abacus	PC
copy	9.98	99.9	38.2	85.5
reverse	10.08	45.5	14.13	78.5
multi.add	4.84	69.23	6.35	41.5
sort	14.55	39.72	21.78	37.7
N*2 multiply	17.06	69.15	18.95	93.7

Figure 3: Methods for Length Generalization on multiple task.

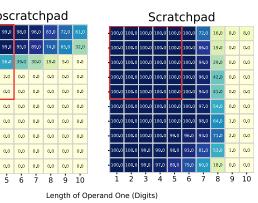


Figure 4: Performance comparison on multiplication.

## Conclusions

**Conclusion.** We show that a plain decoder-only Transformer can achieve strong length generalization on arithmetic tasks through a simple change in data format. Experiments on addition, copy, reverse, multi-operand addition, sorting, and  $N \times 2$  multiplication confirm that ABA consistently improves length generalization over standard formats. It also improves multiplication generalization with a scratchpad.

**Future work.** Future work includes applying ABA-style aligned formatting to broader classes of algorithmic and reasoning tasks where key tokens have a natural alignment structure. Another direction is to study how ABA interacts with larger pretrained language models by reformatting arithmetic or algorithmic prompts without retraining the backbone model.

## Reference

- [1] Sean McLeish, Arpit Bansal et.al. Transformers can do arithmetic with the right embeddings, Nuerips 2024.
- [2] Hanseul Cho, Jaeyoung Cha et.al. Position coupling: Improving length generalization of arithmetic transformers using task structure, Nuerips 2024.
- [3] Nayoung Lee, Kartik Sreenivasan et.al. Teaching arithmetic to small transformers, ICLR 2024.